

R Notebook

This is an R Markdown (<http://rmarkdown.rstudio.com>) Notebook. When you execute code within the notebook, the results appear beneath the code.

```
#initial stuff
rm(list = ls())
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr  1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.1.1      v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
#see all the digits (hopefully)
options(digits = 22)
memory.limit(20000)
```

```
## [1] 20000
```

```
setwd("E:/Duke/Semester 2/Applied Micro/Assignments/A1")
getwd()
```

```
## [1] "E:/Duke/Semester 2/Applied Micro/Assignments/A1"
```

```
#Loading in data
# find all files and then import them separately
dataFiles = list.files("E:/Duke/Semester 2/Applied Micro/Assignments/A1",pattern="*.csv") #all d
ata should be .csv files so this should get all the names (including .csv extension)
# for (i in 1:(length(dataFiles)/2)) assign(dataFiles[i], read.csv(dataFiles[i],colClasses = c
("idmen" = "character"),header = TRUE))
# for (i in 1+(length(dataFiles)/2):length(dataFiles)) assign(dataFiles[i], read.csv(dataFiles
[i],colClasses = c("idind" = "character","idmen" = "character"),header = TRUE))

for (i in 1:length(dataFiles)) assign(dataFiles[i], read.csv(dataFiles[i],colClasses = c("idind"
="character","idmen" = "character"),header = TRUE))
```

[illegible]

```
#got warnings about colClasses sometimes not existing - believe that is because idind is not in
hhDatas

#read all IDs as characters as there were issues with repetition due to R not fully reading values
(some truncation issue)

#Get warnings because the hhData doesn't have idind and I just changed all the idmen and idind to
characters for the issue's discussed in slack
```

Exercise 1

```
#a. Number of Households in 2007
nrow(dathh2007.csv)
```

```
## [1] 10498
```

```
numHH07 <- dathh2007.csv%>%group_by(idmen)%>%summarize(count= n())
sum(numHH07$count)
```

```
## [1] 10498
```

```
#b. Number of households with marital status - Couple with Kids - in 2005
# used table function and just checked when the mstatus was True in 2005
table(dathh2005.csv$mstatus == "Couple, with Kids")
```

```
##
## FALSE TRUE
## 6379 3374
```

```
#c. Number of individuals surveyed in 2008
#
numIND08 <- datind2008.csv%>%group_by(idind)%>%summarize(count= n())
sum(numIND08$count)
```

```
## [1] 25510
```

```
#each row should be unique person so these two, if data is good, should be the same
nrow(datind2008.csv)
```

```
## [1] 25510
```

```
#d. Number of individuals between 25 and 35 in '16
# made a condition that only selected ages between 25 and 35 and chose that subset of the corresponding variable in the right year.
condition <- datind2016.csv$age >= 25 & datind2016.csv$age <= 35
#length gives us the number of people between those ages
length(datind2016.csv$age[condition])
```

```
## [1] 2765
```

```
#e. Cross-Table Gender/Profession in 2009
table(datind2009.csv$gender, datind2009.csv$profession)
```

```
##
##           0  11  12  13  21  22  23  31  33  34  35  37  38  42  43  44  45
## Female  11  30   8  29  63  65   8  68  85 184  50 179  78 258 437   1 153
## Male    19  57  19  78 213 114  48  98 107 142  59 260 368 110 117   2  95
##
##           46  47  48  52  53  54  55  56  62  63  64  65  67  68  69
## Female 410  82  22 782  27 584 353 696  64  35  29  19 147 120  40
## Male   340 429 215 169 182  98 101  74 443 520 246 159 237 177  82
```

```
#profession is a categorical variable so it will be some numbers
```

#f. distribution of wages in 2005 and 2019.

#####

#calculating Gini

for simplicity, I will assume that individuals in each "bucket" earn the same amount

ie if, in decile 5, the value is 500 - I assume everybody in the 10% of all make 500.

for extra simplicity (scaling), will assume 1 person in each bucket (so in theory, we could just scale up all wages but shouldn't change stuff)

#get deciles (10% - 100%). Also showing Lorenz Curve and Equality Curves

#####

GINI <- function(inputData){

deciles <- quantile(inputData,seq(.1,1,by=.00001)) # really small increments so we can accurately get areas and get corresponding values for deciles

#extract data into vector

decile_Vector <- c()

for (i in 1:length(deciles)){

decile_Vector <- c(decile_Vector,deciles[[i]])

}

#find total income (needed for finding % of total income at an index point)

totalIncome <- sum(decile_Vector) #finding total

cumInc <- cumsum(decile_Vector)/totalIncome #find cumulative income at each of the decile levels

#plot for visuals

plot(cumInc, type = "l",lwd = 3, ylab = "Cum Income Share") #"Lorenz" Curve approximation

abline(a=0,b=1/length(deciles), lwd = 3, lty = "dotted") # full equality line

find area between curves ratio

library(pracma) # needed to find area

underLorenz <- trapz(1:length(deciles),cumInc) # way to calculate area using trapezoidal approximation

*area <- length(deciles)*1*.5*

gini <- (area-underLorenz)/area

return(gini)

}


separately

#2005

#cannot do certain calculations with missing data so choosing to omit that. Many 0's but keeping that in as it may be likely that many individuals don't have an income for one reason or another

hist(datind2005.csv\$wage[!is.na(datind2005.csv\$wage & datind2005.csv\$wage != 0)])

am of datind2005.csv\$wage[!is.na(datind2005.csv\$wage & datind2005.csv\$



datind2005.csv\$wage[!is.na(datind2005.csv\$wage & datind2005.csv\$wage != 0)]

```
mean(datind2005.csv$wage[!is.na(datind2005.csv$wage) & datind2005.csv$wage != 0]) # cannot do
mean (returns NA) if we don't omit missing data
```

```
## [1] 22443.02911846829
```

```
sd(datind2005.csv$wage[!is.na(datind2005.csv$wage) & datind2005.csv$wage != 0]) # cannot do me
an (returns NA) if we don't omit missing data
```

```
## [1] 18076.708881794755
```

```
D1 <- quantile(datind2005.csv$wage[!is.na(datind2005.csv$wage)& datind2005.csv$wage != 0],.1)
D1
```

```
## 10%
```

```
## 4547
```

```
D9 <- quantile(datind2005.csv$wage[!is.na(datind2005.csv$wage)& datind2005.csv$wage != 0],.9)
D9 # is 32340.4
```

```
##           90%
## 40452.500000000007
```


```
idr05 <- D9/D1
idr05[[1]]
```

```
## [1] 8.8965251814383119
```

```
#input into created function
inputData05 <- datind2005.csv$wage[!is.na(datind2005.csv$wage) & datind2005.csv$wage != 0]
GINI(inputData05)
```

```
##
## Attaching package: 'pracma'
```

```
## The following object is masked from 'package:purrr':
##
## cross
```




```
## [1] 0.3201291433600677
```

#2019

```
hist(datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0])
```

am of datind2019.csv\$wage[!is.na(datind2019.csv\$wage) & datind2019.csv



datind2019.csv\$wage[!is.na(datind2019.csv\$wage) & datind2019.csv\$wage != 0]

```
mean(datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0])
```

```
## [1] 27578.839302189048
```

```
sd(datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0])
```

```
## [1] 25107.187195539096
```

```
D1 <- quantile(datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0],.1)
D1 # is 0
```

```
## 10%
## 3634.0000000000005
```

```
D9 <- quantile(datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0],.9)
D9 # is 40267
```




```
##           90%
## 50375.600000000006
```

```
idr19 <- D9/D1
idr19[[1]]
```

```
## [1] 13.862300495321959
```


```
inputData19 <- datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0]
GINI(inputData19)
```



```
## [1] 0.33903644704838326
```

```
# together
#combine wage data together into vector
wagesIn05And19 <- c(datind2005.csv$wage[!is.na(datind2005.csv$wage) & datind2005.csv$wage !=0]
,datind2019.csv$wage[!is.na(datind2019.csv$wage) & datind2019.csv$wage !=0])
hist(wagesIn05And19)
```

Histogram of wagesIn05And19



```
#make sure no missing values
# wagesIn05And19 <- wagesIn05And19[!is.na(wagesIn05And19)]
# wagesIn05And19
mean(wagesIn05And19)
```

```
## [1] 25232.617967290786
```

```
sd(wagesIn05And19)
```

```
## [1] 22320.350540064475
```

```
D1 <- quantile(wagesIn05And19,.1)
D1
```

```
## 10%
## 3991
```


```
D9 <- quantile(wagesIn05And19,.9)
D9
```

```
## 90%
## 46219
```

```
idr05And19 <- D9/D1
idr05And19[[1]]
```

```
## [1] 11.580806815334503
```


```
inputData05And19 <- wagesIn05And19
GINI(inputData05And19)
```



```
## [1] 0.33725692054417167
```

```
#g. distributio in age in 2010
hist(datind2010.csv$age)# general
```

Histogram of datind2010.csv\$age



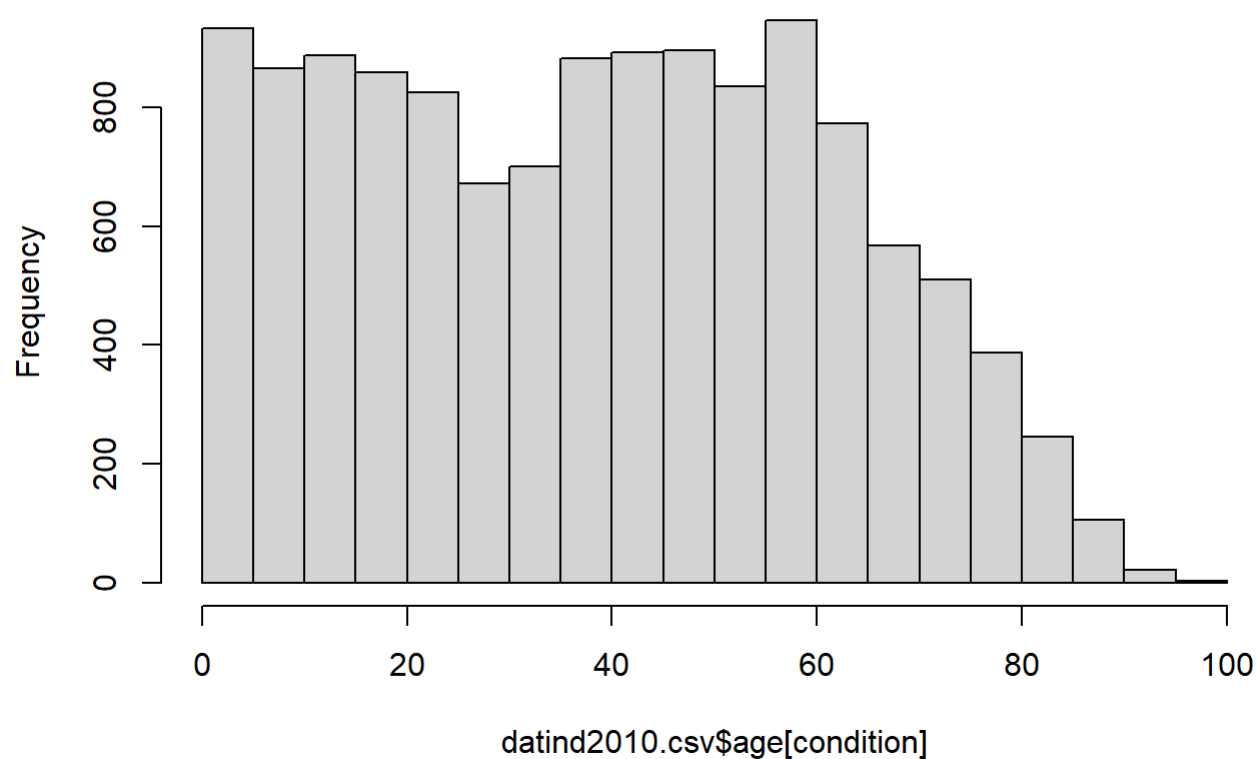
```
summary(datind2010.csv$age)
```

```
##           Min.          1st Qu.          Median          Mean
##  0.0000000000000000  19.000000000000000  40.000000000000000  39.878934077117336
##           3rd Qu.          Max.
##  58.000000000000000 102.000000000000000
```

```
#defining condition so that we look at just Males (to compare to females)
condition <- datind2010.csv$gender == "Male"

#shape of distributions are similar but not identical so there is a difference between men and women
hist(datind2010.csv$age[condition])
```

Histogram of datind2010.csv\$age[condition]




```
summary(datind2010.csv$age[condition])
```

```
##           Min.           1st Qu.           Median           Mean  
## 0.0000000000000000 19.000000000000000 39.000000000000000 38.873623369522768  
##           3rd Qu.           Max.  
## 57.000000000000000 96.000000000000000
```

```
hist(datind2010.csv$age[!condition])
```

Histogram of datind2010.csv\$age[!condition]



```
summary(datind2010.csv$age[!condition])
```

```
##           Min.          1st Qu.          Median          Mean
##  0.000000000000000  20.0000000000000  42.0000000000000  40.816506410256409
##           3rd Qu.          Max.
##  59.0000000000000 102.0000000000000
```

```

#h. number of individuals in Paris in 2011
#creating a vector of house id (assumed all unique even if idind match for some reason)
hhNum <- datind2011.csv$idmen
# hhNum
# dathh2011.csv[998,8] == "Paris" # should be NA
#count occurrences of Paris

countInParis <- 0
for (i in 1:length(hhNum)){
  #find matching household
  index <- match(hhNum[i],dathh2011.csv[,2])
  # if household location is not NA and is Paris, add to count
  if (!is.na(dathh2011.csv[index,8])){
    if(dathh2011.csv[index,8] == "Paris"){
      countInParis <- countInParis + 1
    }
  }
}

countInParis # 3514

```

```
## [1] 3514
```

```

#pipeline group individuals by household number and summarize number of people in each house
hold
anotherWay <- datind2011.csv %>% group_by(idmen) %>% summarize(count = n())

# anotherWay

counter <- 0
for (i in 1:nrow(anotherWay)){
  #extract house id
  hhID <- anotherWay[[i,1]]
  #find indexed match in household data
  index <- match(hhID,dathh2011.csv[,2])
  #if household in Paris, add total number of household people to current value
  if (dathh2011.csv[index,8] == "Paris" & !is.na(dathh2011.csv[index,8])){
    counter <- counter + anotherWay[[i,2]]
  }
}
counter

```

```
## [1] 3514
```

```
#same answer!
```

Exercise 2

```
#a. append all the individual datasets
#a read all the data (previously done) and combined them to form aggregate individual data
#merge individual data
indData <- datind2004.csv
indData <- rbind(indData,datind2005.csv)
indData <- rbind(indData,datind2006.csv)
indData <- rbind(indData,datind2007.csv)
indData <- rbind(indData,datind2008.csv)
indData <- rbind(indData,datind2009.csv)
indData <- rbind(indData,datind2010.csv)
indData <- rbind(indData,datind2011.csv)
indData <- rbind(indData,datind2012.csv)
indData <- rbind(indData,datind2013.csv)
indData <- rbind(indData,datind2014.csv)
indData <- rbind(indData,datind2015.csv)
indData <- rbind(indData,datind2016.csv)
indData <- rbind(indData,datind2017.csv)
indData <- rbind(indData,datind2018.csv)
indData <- rbind(indData,datind2019.csv)

length(unique(indData$idind)) #42868 with numerics. 100160 with characters
```

```
## [1] 100160
```

```
length(unique(indData$idmen)) #41086 with numerics. 41086 with characters
```

```
## [1] 41086
```

```
#b. similar process as above but did it for household data
#merge household data (by row)
hhData <- dathh2004.csv
hhData <- rbind(hhData,dathh2005.csv)
hhData <- rbind(hhData,dathh2006.csv)
hhData <- rbind(hhData,dathh2007.csv)
hhData <- rbind(hhData,dathh2008.csv)
hhData <- rbind(hhData,dathh2009.csv)
hhData <- rbind(hhData,dathh2010.csv)
hhData <- rbind(hhData,dathh2011.csv)
hhData <- rbind(hhData,dathh2012.csv)
hhData <- rbind(hhData,dathh2013.csv)
hhData <- rbind(hhData,dathh2014.csv)
hhData <- rbind(hhData,dathh2015.csv)
hhData <- rbind(hhData,dathh2016.csv)
hhData <- rbind(hhData,dathh2017.csv)
hhData <- rbind(hhData,dathh2018.csv)
hhData <- rbind(hhData,dathh2019.csv)

length(unique(hhData$idmen)) #41084 with numerics and characters
```



```
## [1] 41084
```

```
#c. List all the variables present in both  
# colnames(indData)  
# colnames(hhData)  
intersect( colnames(indData),  
           colnames(hhData))
```

```
## [1] "X"      "idmen" "year"
```

```
#matching variables: X, idmen, year
```

```
#d. merge the appended data sets  
# creating subsets of data that have matching variables () idmen, and year) - omitted X as it was  
just a byproduct of reading in the data and would just mess up stuff  
  
aggData <- merge(hhData,indData,by = c("idmen","year"))  
nrow(aggData)
```

```
## [1] 413501
```

```
#e. Number of HH with more than 4 people  
#group by HH and year and then find how many have more than 4 people associated with it  
aggDatae <- aggData %>% group_by(idmen,year) %>% summarise(count = n())
```

```
## `summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.
```

```
# aggDatae  
length(which(aggDatae[, "count"] > 4))
```

```
## [1] 12436
```

```
#f. Number of HH in which at least one member is UE  
#pipeline group by idmen and year and see if ANY individual, in the household for a given year,  
is Unemployed by using %in%  
# will return TRUE if at least 1 person in Unemployed  
aggDataf <- aggData %>% group_by(idmen,year) %>% summarise(UE = "Unemployed" %in% empstat)
```

```
## `summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.
```

```
# aggDataf  
  
length(which(aggDataf$UE == TRUE))
```

```
## [1] 17241
```

#g. Number of HH in which at least 2 have the same profession

group by household and year and filter out professions that are missing or blank and then see if there are any Duplicates (anyDuplicated will return index of repeat if there is a duplicate and 0 otherwise)

```
aggDatag <- aggData %>% group_by(idmen,year) %>% filter(!is.na(profession), profession != "") %>% summarise(sameProf = anyDuplicated(profession))
```

`summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.

```
length(which(aggDatag$sameProf != 0))
```

```
## [1] 7586
```

#h. Num of individuals that are from Couple with Kids

each row in aggData should be unique individual so just go through rows and check

```
countHHwithKids <- 0
```

each row should represent an individual in aggregate data so simply checking if that person is in a household "Couple, with Kids"

```
for (i in 1:nrow(aggData)){
  if(aggData$mstatus[i] == "Couple, with Kids" & !is.na(aggData$mstatus[i])){
    countHHwithKids <- countHHwithKids + 1
  }
}
countHHwithKids
```

```
## [1] 209382
```

quicker way

```
length(which(aggData$mstatus == "Couple, with Kids"))
```

```
## [1] 209382
```

```

#i. number of individuals in Paris
#similar process as h - go through row of aggData and see if person in Paris. However, with ques
tion wording, I was unsure if it meant how many individuals were in Paris at some point in time
in their survey responses so I did both

#if meant individual in a year - by row
countInParis <- 0
#each row should represent individual so simply check if that person is in "Paris"
for (i in 1:nrow(aggData)){

  if (!is.na(aggData$location[i])){
    if (aggData$location[i] == "Paris"){
      countInParis <- countInParis + 1
    }
  }
}
countInParis

```

```
## [1] 51904
```

```

#another way
length(which(aggData$location == "Paris"))

```

```
## [1] 51904
```

```

# #if mean unique individuals ever in Paris
# uniqueIdind <- unique(aggData$idind)
# aggData_i = aggData %>% group_by(idind) %>% summarize(fromParis = "Paris" %in% location)
# length(which(aggData_i$fromParis == TRUE))

```

```

# j. idmen of largest household(s)
#group by household and year and measure size of group
aggData_j <- aggData %>% group_by(idmen, year) %>% summarise(count = n())

```

```
## `summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.
```

```

#find largest households in a given year
bigFam <- max(aggData_j[,3])
#find index position of those
which(aggData_j[,3] == bigFam) #there were 2

```

```
## [1] 69653 107200
```

```
#find indmen and year
idmen1Max <- aggDataj[[which(aggDataj[,3] == 14)[1],1]]
idmen2Max <- aggDataj[[which(aggDataj[,3] == 14)[2],1]]
idmen1MaxYear <- aggDataj[[which(aggDataj[,3] == 14)[1],2]]
idmen2MaxYear <- aggDataj[[which(aggDataj[,3] == 14)[2],2]]

#reporting them
idmen1Max
```

```
## [1] "2207811124040100"
```

```
idmen1MaxYear
```

```
## [1] 2007
```

```
idmen2Max
```

```
## [1] "2510263102990100"
```

```
idmen2MaxYear
```

```
## [1] 2010
```

```
#k. Number of households present in 2010 and 2011
nrow(dathh2010.csv) + nrow(dathh2011.csv)
```

```
## [1] 22408
```

```
# IF QUESTION MEANT HOW MANY HOUSEHOLDS WERE BOTH IN 2010 and 2011 (ie idmen in both 2010 and 2011) but nothing about respondents
```

```
aggData3 <- aggData %>% group_by(idmen) %>% summarise(present3 = ((2010 %in% year & 2011 %in% year)))
# aggData3
```

```
length(which(aggData3$present3 == TRUE))
```

```
## [1] 8984
```

Exercise 3

#a. Find out each year household enters and exit the panel - report distribution of time spent for each household

found total unique households

```
uniqueIdmen <- unique(aggData$idmen) #getting unique household IDs
length(uniqueIdmen) #41084 unique households
```

```
## [1] 41084
```

mutate data by household and put entry, exit

```
hhEntryExit = aggData %>% group_by(idmen) %>% filter(!is.na(year)) %>% mutate(entryYear = min(year), exitYear = max(year))
```

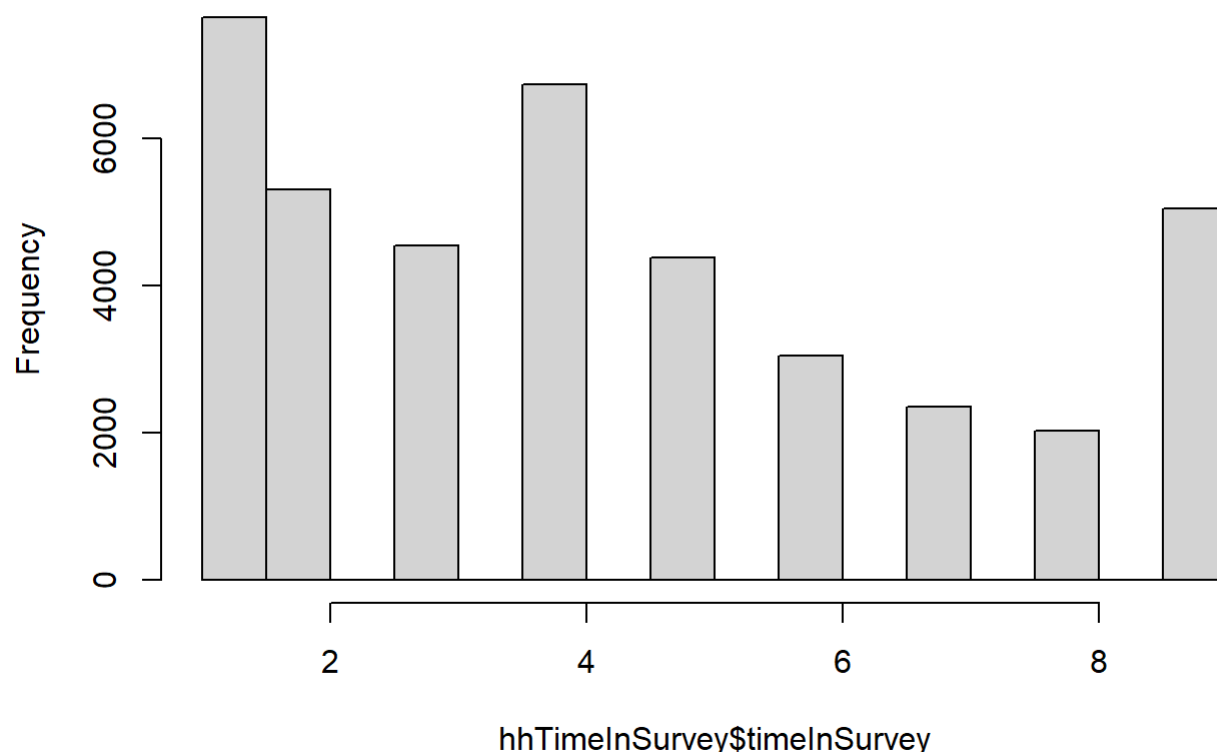
by household, find their time in survey by taking the difference of their last and first year and adding 1

```
hhTimeInSurvey = hhEntryExit %>% group_by(idmen) %>% summarise(timeInSurvey = max(year) - min(year) + 1)
```

not getting unique households but individual level time spent if i don't use hhTimeInSurvey

```
hist(hhTimeInSurvey$timeInSurvey)
```

Histogram of hhTimeInSurvey\$timeInSurvey



```
summary(hhTimeInSurvey$timeInSurvey)
```

```
##              Min.             1st Qu.             Median             Mean
## 1.0000000000000000 2.0000000000000000 4.0000000000000000 4.3123357024632458
##              3rd Qu.             Max.
## 6.0000000000000000 9.0000000000000000
```

```
#
# hhTimeInSurvey = aggData %>% group_by(idmen) %>% summarize(timeInSurvey = max(year) - min(year)+1)
# hist(hhTimeInSurvey$timeInSurvey)
# summary(hhTimeInSurvey$timeInSurvey)
#
```

```
#b. based on datent, check where household moved into current dwelling at time of year

# breaking into 2 parts since, for me, it seems more logical

#part 1 - Looking at HH data and checking if each house had year == datent
hhDataByYearb <- hhData %>% group_by(idmen,year) %>% mutate(movedAtYearOfSurvey = year == date
nt) # group by year and will be TRUE if year is same as datent

#reporting first 10 rows of relevant variables
View(hhDataByYearb[1:10,c("idmen","year","datent","movedAtYearOfSurvey")])

nrow(hhDataByYearb)
```

```
## [1] 173851
```

```
length(which(hhDataByYearb$movedAtYearOfSurvey == TRUE))
```

```
## [1] 4776
```

#part 2

group by year and ID migration

```
method1ByYearMigration <- aggData %>% group_by(year) %>% filter(!is.na(year),!is.na(datent)) %>%
summarise(migration = length(which(year == datent)))
```


```
method1ByYearTotal <- aggData %>% group_by(year) %>% filter(!is.na(year),!is.na(datent)) %>% sum
marise(count = n())
```

#find proportions

```
shareofPopByYear <- method1ByYearMigration[,2]/method1ByYearTotal[,2]
```

plotting

```
shareOfPopdf <- data.frame(c(04,05,06,07,08,09,10,11,12,13,14,15,16,17,18,19),shareofPopByYear)
colnames(shareOfPopdf) <- c("year","Proportion")
ggplot(shareOfPopdf, aes(x= year,y = Proportion))+
  geom_bar(stat="identity")
```



```
shareOfPopM1 <- shareofPopByYear
```

*#b. similar thing but using myear and move. myear available until 2014 and move available after.
Assuming move == 2 means moved (instead of 1)*

```
# first part - whether households migrated that year or not
hhDataByYearc <- hhData %>% group_by(idmen,year) %>% mutate(
  migration = if (year <= 2014){
    migration = year == myear # TRUE if year is equal to the move year
  }
  else if(year > 2014){
    migration = move == 2 # TRUE if they moved since last year (ie they moved this year)
  }
)
```

```
View(hhDataByYearc[1:10,c("idmen","year","myear","move","migration")])
```

#part2 - same process as earlier, just different checks

```
method2ByYearMigration <- aggData %>% group_by(year) %>% summarise(migration = length(which(
  if (year <= 2014){
    migration = year == myear # TRUE if year is equal to the move year
  }
  else if(year > 2014){
    migration = move == 2 # TRUE if they moved since last year (ie they moved this year)
  })))
```



```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year > 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year > 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year > 2014) {: the condition has length > 1 and only the first  
## element will be used
```


```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year > 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year <= 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
## Warning in if (year > 2014) {: the condition has length > 1 and only the first  
## element will be used
```

```
method2ByYearTotal <- aggData %>% group_by(year) %>% summarise(count = n())  
  
shareofPopByYear <- method2ByYearMigration[,2]/method2ByYearTotal[,2]  
shareOfPopM2 <- shareofPopByYear  
  
shareOfPopdf <- data.frame(c(04,05,06,07,08,09,10,11,12,13,14,15,16,17,18,19),shareofPopByYear)  
colnames(shareOfPopdf) <- c("year","Proportion")  
ggplot(shareOfPopdf, aes(x= year,y = Proportion))+  
  geom_bar(stat="identity")
```



```
nrow(hhDataByYearc)
```

```
## [1] 173851
```

```
length(which(hhDataByYearc$migration == TRUE))
```


```
## [1] 6071
```

```
#d. merge the data
```

```
method1 <- rep("Method1",16)
method2 <- rep("Method2",16)
method <- c(method1,method2)
shareOfProportion <- rbind(shareOfPopM1,shareOfPopM2)

#create data frame
combinedData <- data.frame(c(04,05,06,07,08,09,10,11,12,13,14,15,16,17,18,19),method,shareOfPr
oportion)
# View(combinedData)
colnames(combinedData) <- c("year","Method","Proportion")

ggplot(combinedData, aes(x=year, y = Proportion, fill = Method)) +
  geom_bar(stat = "identity",position = "dodge")
```



```
length(which(is.na(aggData[,2]))) #none missing
```

```
## [1] 0
```

```
length(which(is.na(aggData[,4])))
```

```
## [1] 245
```

```
#245 missing
```

```
in04 <- which(aggData[,2] == 2004)

in05 <- which(aggData[,2] == 2005)

in06 <- which(aggData[,2] == 2006)

in07 <- which(aggData[,2] == 2007)

in08 <- which(aggData[,2] == 2008)

in09 <- which(aggData[,2] == 2009)

in10 <- which(aggData[,2] == 2010)

in11 <- which(aggData[,2] == 2011)

in12 <- which(aggData[,2] == 2012)

in13 <- which(aggData[,2] == 2013)

in14 <- which(aggData[,2] == 2014)

in15 <- which(aggData[,2] == 2015)

in16 <- which(aggData[,2] == 2016)

in17 <- which(aggData[,2] == 2017)

in18 <- which(aggData[,2] == 2018)

in19 <- which(aggData[,2] == 2019)
# below only work if you uncomment inXY from above
until14 <- c(in04,in05,in06,in07,in08,in09,in10,in11,in12,in13,in14)
after14 <- c(in15,in16,in17,in18,in19)
length(which(is.na(aggData[until14,5])))
```

```
## [1] 6446
```

```
#6446 missing
```

```
length(which(is.na(aggData[after14,7])))
```

```
## [1] 25250
```

```
#25250 missing
```

#e. Find out how many households had at least one family member change profession/empstat

first issue is determining which criteria to use to do this - question said migrate and that was associated with method2 as opposed to method1(which used moved). That is why I chose to use that method. If time permits, I will also try to do it with method1 as I have previously said it is better (more data)

#using method2

```
method2Migration <- hhData %>% group_by(idmen,year) %>% summarise(  
  migration = if (year <= 2014){  
    migration = year == myear # TRUE if year is equal to the move year  
  }  
  else if(year > 2014){  
    migration = move == 2} # TRUE if they moved since last year (ie they moved this year)
```

`summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.

```
#individuals in houses that have migrated (according to method2)
```

```
newAggData <- merge(aggData,method2Migration, by = c("idmen","year"))
#add column of 0s to newAggData to track job movements
Change <- rep(0,nrow(newAggData))
newAggData <- cbind(newAggData,Change) #add column of zeros at the end which will then be used
for tracking job changes
migratedHouseData <- newAggData[ which(newAggData$migration == TRUE),]
```

```
#extract unique people
```

```
uniqueIdind <- unique(migratedHouseData$idind)
# Length(uniqueIdind)
```

```
#go through list
```

```
for (i in 1:length(uniqueIdind)){
  ID <- uniqueIdind[i]
  #find indices corresponding to them
  index <- which(migratedHouseData[, "idind"] == ID)
```

```
#check if empstat or profesison changed
```

```
for (j in 1:length(index)){
  if (j > 1){
```

```
    #get index for years
    lastYearIndex <- index[j-1]
    thisYearIndex <- index[j]
```

```
    if(!is.na(migratedHouseData[lastYearIndex,"profession"]) &
       !is.na(migratedHouseData[thisYearIndex,"profession"])){
#checked if not missing professions between this year and last year
      #check if not blank both years and if they are not the same (ie change occurred)
      if((migratedHouseData[lastYearIndex,"profession"] != "" &
          migratedHouseData[thisYearIndex,"profession"] != "") &
          (migratedHouseData[lastYearIndex,"profession"] !=
            migratedHouseData[thisYearIndex,"profession"])){

        migratedHouseData[thisYearIndex,"Change"] <- 1
      }
    }
  }
```

```
# check for change in empstat if profession doesnt change
```

```
if (!is.na(migratedHouseData[lastYearIndex,"empstat"]) &
    !is.na(migratedHouseData[thisYearIndex,"empstat"])){
  if (migratedHouseData[lastYearIndex,"empstat"] != migratedHouseData[thisYearIndex,"emp
```

```
stat"]){
  migratedHouseData[thisYearIndex,"Change"] <- 1
}
}

}

}
}
# group and sum up changes (no changes in profession/employment would equal 0)

# sum will be 0 if no change in that year
jobChangeAndMove <- migratedHouseData %>% group_by(idmen,year) %>% summarise(jobChange = sum(C
hange) )
```

```
## `summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.
```

```
length(which(jobChangeAndMove[,3] > 0))
```

```
## [1] 518
```

```
jobChangeByYear2 <- jobChangeAndMove %>% group_by(year) %>%
summarise(changeInProfOrEmp = length(which(jobChange!= 0)))
```

```
jobChangeByYear2
```

year <int>	changeInProfOrEmp <int>
2004	0
2005	4
2006	12
2007	12
2008	28
2009	27
2010	21
2011	28
2012	34
2013	38

1-10 of 16 rows

Previous **1** 2 Next

```
length(which(jobChangeAndMove[,3] > 0))
```



```
## [1] 518
```

```
#####  
#using method1  
method1Migration <- hhData %>% group_by(idmen,year) %>% summarise(  
  migration = year == datent) # TRUE if they moved since last year (ie they moved this year)
```

```
## `summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.
```

```
#individuals in houses that have migrated (according to method2)
```

```
newAggData <- merge(aggData,method1Migration, by = c("idmen","year"))
#add column of 0s to newAggData to track job movements
Change <- rep(0,nrow(newAggData))
newAggData <- cbind(newAggData,Change) #add column of zeros at the end which will then be used
for tracking job changes
migratedHouseData <- newAggData[ which(newAggData$migration == TRUE),]
```

```
#extract unique people
uniqueIdind <- unique(migratedHouseData$idind)
# Length(uniqueIdind)

#go through list
for (i in 1:length(uniqueIdind)){
  ID <- uniqueIdind[i]
  #find indices corresponding to them
  index <- which(migratedHouseData[, "idind"] == ID)
```

```
#check if empstat or profesison changed
```

```
for (j in 1:length(index)){
  if (j > 1){
```

```
    #get index for years
    lastYearIndex <- index[j-1]
    thisYearIndex <- index[j]
```

```
    if(!is.na(migratedHouseData[lastYearIndex,"profession"]) &
       !is.na(migratedHouseData[thisYearIndex,"profession"])){
```

```
      #checked if not missing professions between this year and last year
      #check if not blank both years and if they are not the same (ie change occurred)
```

```
      if((migratedHouseData[lastYearIndex,"profession"] != "" &
          migratedHouseData[thisYearIndex,"profession"] != "") &
          (migratedHouseData[lastYearIndex,"profession"] !=
           migratedHouseData[thisYearIndex,"profession"])){
```

```
        migratedHouseData[thisYearIndex,"Change"] <- 1
      }
    }
```

```
# check for change in empstat if profession doesnt change
```

```
    if (!is.na(migratedHouseData[lastYearIndex,"empstat"]) &
```

```

      !is.na(migratedHouseData[thisYearIndex,"empstat"])){
    if (migratedHouseData[lastYearIndex,"empstat"] != migratedHouseData[thisYearIndex,"emp
stat"]){
      migratedHouseData[thisYearIndex,"Change"] <- 1
    }
  }

}

}

}
# group and sum up changes (no changes in profession/employment would equal 0)

# sum will be 0 if no change in that year
jobChangeAndMove <- migratedHouseData %>% group_by(idmen,year) %>% summarise(jobChange = sum(C
hange) )

```

`summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.

```
length(which(jobChangeAndMove[,3] > 0))
```

```
## [1] 312
```

```

jobChangeByYear1 <- jobChangeAndMove %>% group_by(year) %>%
  summarise(changeInProfOrEmp = length(which(jobChange!= 0)))
jobChangeByYear1

```

year <int>	changeInProfOrEmp <int>
2004	1
2005	5
2006	14
2007	13
2008	24
2009	25
2010	22
2011	30
2012	34
2013	40

1-10 of 16 rows

Previous **1** 2 Next

```
length(which(jobChangeAndMove[,3] > 0))
```

```
## [1] 312
```

```
#####
```

```
#using either
methodMergedMigration <- hhData %>% group_by(idmen,year) %>% summarise(
  migration = if (year <= 2014){
    migration = year == myear # TRUE if year is equal to the move year
  }
  else if(year > 2014){
    migration = move == 2}
  else {migration = year == datent}) # TRUE if they moved since last year (ie they moved this year)
```

```
## `summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.
```

```
#individuals in houses that have migrated (according to method2)
```

```
newAggData <- merge(aggData,methodMergedMigration, by = c("idmen","year"))
#add column of 0s to newAggData to track job movements
Change <- rep(0,nrow(newAggData))
newAggData <- cbind(newAggData,Change) #add column of zeros at the end which will then be used
for tracking job changes
migratedHouseData <- newAggData[ which(newAggData$migration == TRUE),]

#extract unique people
uniqueIdind <- unique(migratedHouseData$idind)
# length(uniqueIdind)

#go through list
for (i in 1:length(uniqueIdind)){
  ID <- uniqueIdind[i]
  #find indices corresponding to them
  index <- which(migratedHouseData[, "idind"] == ID)

  #check if empstat or profesison changed

  for (j in 1:length(index)){
    if (j > 1){

      #get index for years
      lastYearIndex <- index[j-1]
      thisYearIndex <- index[j]

      if(!is.na(migratedHouseData[lastYearIndex,"profession"]) &
        !is.na(migratedHouseData[thisYearIndex,"profession"])){

        #checked if not missing professions between this year and last year
        #check if not blank both years and if they are not the same (ie change occurred)
        if((migratedHouseData[lastYearIndex,"profession"] != "" &
          migratedHouseData[thisYearIndex,"profession"] != "") &
          (migratedHouseData[lastYearIndex,"profession"] !=
            migratedHouseData[thisYearIndex,"profession"])){

          migratedHouseData[thisYearIndex,"Change"] <- 1
        }
      }

      # check for change in empstat if profession doesnt change

      if (!is.na(migratedHouseData[lastYearIndex,"empstat"]) &
        !is.na(migratedHouseData[thisYearIndex,"empstat"])){
        if (migratedHouseData[lastYearIndex,"empstat"] != migratedHouseData[thisYearIndex,"emp
stat"]){
```

```

      migratedHouseData[thisYearIndex,"Change"] <- 1
    }
  }

}

}

# group and sum up changes (no changes in profession/employment would equal 0)

# sum will be 0 if no change in that year
jobChangeAndMove <- migratedHouseData %>% group_by(idmen,year) %>% summarise(jobChange = sum(C
change))

```

`summarise()` has grouped output by 'idmen'. You can override using the `.groups` argument.

```
# jobChangeAndMove
```

```

jobChangeByYearBoth <- jobChangeAndMove %>% group_by(year) %>%
  summarise(changeInProfOrEmp = length(which(jobChange!= 0)))
jobChangeByYearBoth

```

year <int>	changeInProfOrEmp <int>
2004	0
2005	4
2006	12
2007	12
2008	28
2009	27
2010	21
2011	28
2012	34
2013	38

1-10 of 16 rows

Previous **1** 2 Next

```
length(which(jobChangeAndMove[,3] > 0))
```

```
## [1] 518
```

```
View(jobChangeByYear2)
View(jobChangeByYear1)
View(jobChangeByYearBoth)
```

Exercise 4

```
#4
minMaxAttrition <- aggData %>% group_by(idind) %>% summarise(entry = min(year),exit = max(year))

(minMaxAttrition)
```

idind <chr>	entry <int>	exit <int>
1120001001293010001	2004	2004
1120001004058010001	2004	2005
1120001004058010002	2004	2005
1120001006663010001	2004	2005
1120001006663010002	2004	2005
1120001008245010001	2004	2005
1120001008644010001	2004	2005
1120001008644010002	2004	2005
1120001010299010001	2004	2005
1120001010299010002	2004	2005
1-10 of 10,000 rows		
Previous 1 2 3 4 5 6 ... 1000 Next		

```

# now we have entry and exit data for each individual

totalInd <- c()
attritionPerYear <- c()
for (i in 2004:2019){
  year <- i
  count <- 0
  attrition <- 0
  # count how many people in the survey at that time (ie this year is within their entry exit ye
ars)
  count <- length(which(minMaxAttrition$entry <= i & i <= minMaxAttrition$exit))
#find how many people are leaving (ie this year is their exit year)
  attrition <- length(which(minMaxAttrition$exit == i))

  #attach to vectors so we get these values for each year
  totalInd <- c(totalInd,count)
  attritionPerYear <- c(attritionPerYear,attrition)
}

# totalInd
# attritionPerYear
propAttrition <- attritionPerYear/totalInd
years <- seq(2004,2019)
# years
propAttritionTable <- cbind(years,propAttrition)
(propAttritionTable)

```

```

##      years      propAttrition
## [1,] 2004 0.11596820809248555
## [2,] 2005 0.18002351128947261
## [3,] 2006 0.16151297625621203
## [4,] 2007 0.20907050184529924
## [5,] 2008 0.18760132787771172
## [6,] 2009 0.16840324503056633
## [7,] 2010 0.17346976744186046
## [8,] 2011 0.15602553870710295
## [9,] 2012 0.21096051856216852
## [10,] 2013 0.18614624218636430
## [11,] 2014 0.18920394569213408
## [12,] 2015 0.18961038961038962
## [13,] 2016 0.21249587261987746
## [14,] 2017 0.20849655801924466
## [15,] 2018 0.23640940917315711
## [16,] 2019 1.0000000000000000

```